

## SysEng 5212 /EE 5370

### Introduction to Neural Networks and Applications

#### *Week 6 : Backpropagation II and Practical Implementation Considerations*

**Cihan H Dagli, PhD**

*Professor of Engineering Management and Systems Engineering  
Professor of Electrical and Computer Engineering  
Founder and Director of Systems Engineering Graduate Program*

[dagli@mst.edu](mailto:dagli@mst.edu)

MISSOURI UNIVERSITY OF SCIENCE AND TECHNOLOGY  
Rolla, Missouri, U.S.A.

# Volunteers Tally

- Murat Aslan
- Tatiana Cardona Sepulveda
- Prince Codjoe
- Xiongming Dai
- Jeffrey Dierker

# Volunteers Tally

- Venkata Sai Abhishek Dwivadula
- Brian Guenther
- Anthony Guertin
- Timothy Guertin
- Seth Kitchen

# Volunteers Tally

- Gregory Leach
- Yu Li
- John Nganga
- Igor Povarich
- Jack Savage

# Volunteers Tally

- William Symolon
- Wayne Viers III
- Tao Wang
- Kari Ward
- Julia White
- Jun Xu

# Lecture outline

- Midterm exam is on 3/6/2018
  - Open-book + open notes + open Internet, bring your laptops ( Don't forget to credit your sources by giving references)
  - Two and half hours from 4:00-6:30 PM US CST
  - Both on and off campus students will download the exam at the same time from Canvas
  - Turn in the completed exam by uploading it on CANVAS as zip file (on campus students can submit non-electronic version of their answers to Deepak in class)
  - Must sign and attach honor code -exam will not be graded without it

方向传播算法

# Today

## Backprop Review

### Practical Issues in Standard Backprop

- I. Weight initialization
- II. Generalization
- III. Network configuration
- IV. Universal approximation
- V. Curse of dimensionality
- VI. Independent validation
- VII. Bias-variance dilemma
- VIII. Improving convergence speed

# Standard Backpropagation Review

**Step 1:** Initialize the network synaptic weights to small random Values

**Step 2:** Select a random input/output training sample, present it to the network and calculate the network response

**Step 3:** Compute the error in the network response by comparing it to the desired response

MATLAB:  $y = \text{purelin}(v)$



# Backpropagation Review

**Step 4:** Use the output error to calculate all local error gradients  $\delta_j(n)$

- If  $j$  is an output neuron,

$$\delta_j(n) = e_j(n) \phi'_j(v_j(n))$$

- If  $j$  is a hidden neuron,

$$\delta_j(n) = \phi'_j(v_j(n)) \sum_k \delta_k(n) w_{kj}(n)$$

where,

- $e_j(n)$  is the output error
- $\phi'_j(v_j(n))$  is the derivative of the activation function
- $w_{kj}(n)$  is the weight to be updated

**Step 5:** Update the network weights using the weight update or delta rule

$$w_{ji}(n+1) = w_{ji}(n) + \eta \delta_j(n) y_i(n)$$

where  $y_i(n) = x(n)$  for layer 1.

**Step 6:** Continue steps 2 through 5, until the network reaches a desired level of performance

## Practical Considerations in Implementing Backpropagation Learning

1. Initialization of synaptic weights
  - a) Nguyen Widrow initialization algorithm
2. Generalization
  - a) Overfitting
  - b) Network configuration
  - c) Training sample size
3. Universal approximation and network configuration
4. Curse of dimensionality
5. Independent validation
6. Early stopping
7. Bias-variance dilemma
8. Complexity regularization
9. Speed of convergence



# Initializing Network Weights

- Avoid very large and very small values of initial weights.
- Weight initialization heuristic:  $\left[-\frac{0.5}{m_1}, \frac{0.5}{m_1}\right]$  where  $m_1$  is the number of neurons in the first hidden layer
- Alternative weight update method for MLPs with one hidden layer: Nguyen-Widrow's initialization algorithm
  - significantly improves speed of network training



# Nguyen Widrow Initialization Algorithm

Define,  $m_0$  = number of input nodes;  $m_1$  = number of neurons in the hidden layer  
 $\gamma$  = scaling factor

**Step 1:** Compute the scaling factor as,

$$\gamma = 0.7^{m_0} \sqrt{m_1}$$

**Step 2:** Initialize the weights  $w_{ij}$  as uniformly distributed random numbers between -0.5 and 0.5

**Step 3:** Reinitialize weights using,

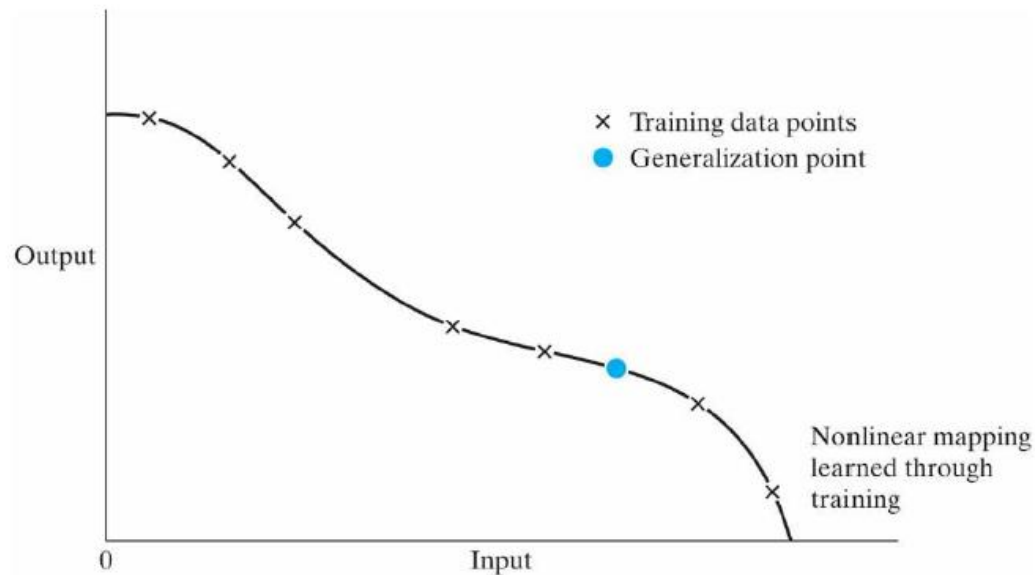
$$w_{ji} = \gamma \frac{w_{ji}}{\sqrt{\sum_{i=1} m_1 w_{ji}^2}}$$

**Step 4:** For the  $j$ th neuron in the hidden layer, set the bias to be a random number between  $w_{ij}$  and  $-w_{ij}$

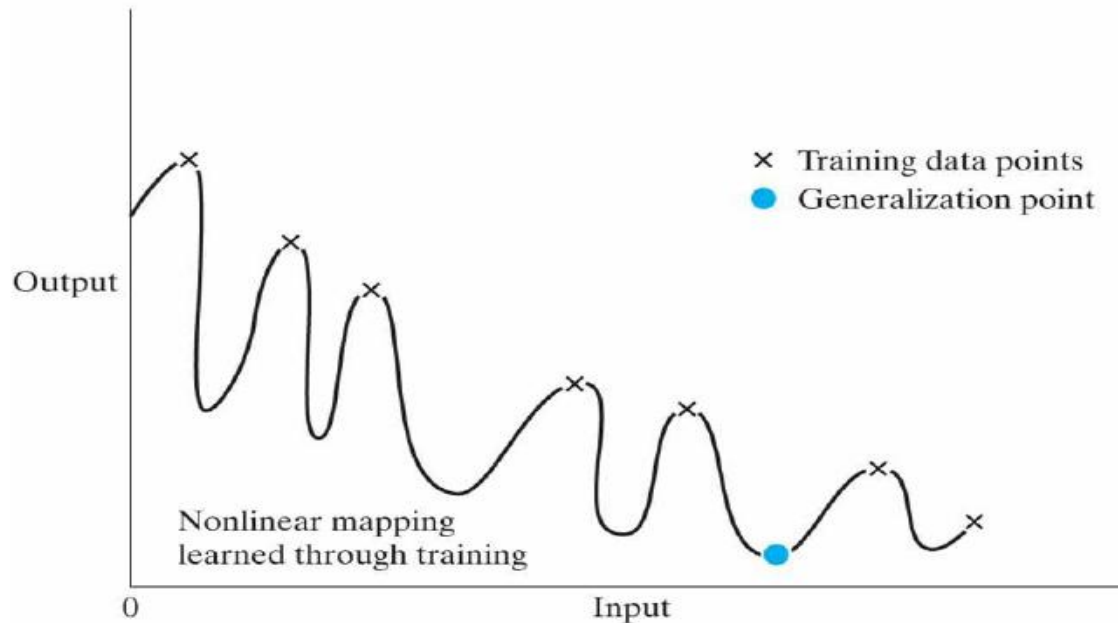


# Generalization

- Training: Encoding input output patterns into the nodes in the form of the synaptic weights
- Generalizing: Correctly computing the response to a never-before-seen input



# Overfitting



If the data contains noise, the network will identify features that don't exist in the function being mapped. Network memorizes the data and loses ability to generalize!

## More on Overfitting

- Overtrained network acts like a "look-up table"  
查出
- Input-output mapping computed by the network is not smooth
- What makes a good mapping?
  - Occam's Razor applies
  - The simplest function mapping the inputs to the outputs is often the best
- The smoothest function is usually the simplest function.





## Factors Affecting Generalization

Good generalization depends on three factors,

- Size of the training sample
- Architecture of the neural network
- Complexity of the problem

We can look at the problem from two perspectives,

- Fix the architecture, find a training sample set that will achieve good generalization
- Fix the training sample, determine the best architecture for generalizing with the available data



# Sufficient Training-Sample Size

In practice, we can achieve good generalization by selecting a training set that satisfies the condition,  $N = O\left(\frac{W}{\varepsilon}\right)$

where,  $W$  is the number of free parameters (weights and biases),  $\varepsilon$  is the fraction of classification errors permitted, and  $O(.)$  is the order of the error fraction.

As a rule of thumb, for good generalization, the number of training examples should be greater than the ratio of the total number of free parameters ( $m_0 + m_1$ ) to the mean-square error estimate.



# Impact of Network Configuration on Performance

Network configuration is determined by,

- Number of hidden layers
- Number of neurons in each hidden layer
- Type of activation functions used

An overdesigned architecture will overfit data; network cannot generalize.

Choice of activation function has little impact on performance as long as it is nonlinear

- Output layer functions may be nonlinear or linear

## More on Network Configuration

- Determining a sufficient number of hidden neurons is frequently done by trial and error
- MLP with one hidden layer acts as a universal approximator
  - Use of more than two hidden layers is rarely justifiable in terms of performance

## Multilayer Perceptron as a Universal Approximator

A multilayer perceptron trained with backpropagation performs a nonlinear mapping of input to the output.

If  $m_0$  is the number of input variables and  $m_L$  is the number of output neurons,

- $m_0$  -dimensional input space  $\rightarrow$   $m_L$  -dimensional output space

What is the minimum number of hidden layers necessary to successfully approximate any given nonlinear input-output mapping?

- The *universal approximation theorem* answers this question.



# Universal Approximation Theorem

Let  $\phi(\cdot)$  be a nonconstant, bounded, and monotone-increasing continuous function. Let  $I_{m_0}$  denote the  $m_0$ -dimensional unit-hypercube  $[0, 1]^{m_0}$ . The space of continuous functions on  $I_{m_0}$  is denoted by  $C(I_{m_0})$ . Then, given any function  $f \in C(I_{m_0})$  and  $\epsilon > 0$ , there exists an integer  $m_1$  and set of real constants  $\alpha_j$ ,  $b_j$ , and  $w_{ji}$  where  $i = 1, \dots, m_0$  and  $j = 1, \dots, m_1$  such that we may define,

$$F(x_1, \dots, x_{m_0}) = \sum_{j=1}^{m_1} \alpha_j \phi\left(\sum_{i=1}^{m_0} w_{ji} x_i + b_j\right)$$

as an approximate realization of the function  $f(\cdot)$ ; that is,

$$|F(x_1, \dots, x_{m_0}) - f(x_1, \dots, x_{m_0})| < \epsilon$$

for all  $x_1, \dots, x_{m_0}$  that lie in the input space.





# Applying the Universal Approximation Theorem to the MLP

$\Phi(.)$ : a nonconstant, bounded, and monotone-increasing continuous function. The sigmoidal nonlinearities as activation functions satisfy this condition

- $m_0$  -dimensional input space: The network input has  $m_0$  nodes denotes by  $x_1, \dots, x_{m_0}$

integers  $m_1$  and set of real constants  $\alpha_j$ ,  $b_i$ , and  $w_{ij}$ : The hidden layer has  $m_1$  neurons;  $\alpha_j$ ,  $b_i$ , and  $w_{ij}$  are the output weights, hidden layer bias and weights respectively

# MLP with Single Hidden Layer is a Universal Approximator

Stated simply,

The theorem states that a single hidden layer is sufficient for a multilayer perceptron to compute an approximation represented by a given input-output mapping.

The theorem guarantees approximation, not exact representation or optimal implementation!



# Practical Considerations In Designing a Multilayer Network

- The universal approximation theorem is not constructive; does not tell us how to determine the most optimal network configuration
- Neurons in a single hidden layer interact with each other
- This makes it difficult to improve the approximation at one point without worsening it at some other point
- Research provides some justification for the use of two hidden layers,
  - first hidden layer extracts local features; creates a new feature space
  - Second hidden layer extracts global features partitions new feature space into desired classes
  - Two layers make the approximation much more manageable



# Curse of Dimensionality

- Problem complexity increases with increase in number of input dimensions
- Functions with higher-dimensional inputs are much more complex than functions with lower dimensional inputs; Are harder to learn
- The complexity is much more expensive to model
- Modeling outcomes are more likely to be inaccurate
- For accurate outcomes the number of input samples  $N$  needs to be,

$$N \gg m_0^2$$


where  $m_0$  is the dimensionality of the input.

- Otherwise, the model is too complicated for the data which leads to overfitting!



# Curse of Dimensionality: Number of input features

Consider a pattern recognition problem with two classes. To start with train the network with just one input feature




Sample density is 10 samples per unit length.




# Increase the Number of Features

- Increasing the number of input features improves discrimination but decreases sample density



- To get the same sample density as one feature,  $10^2$  samples are needed!

With three features,  $10^3$  samples will be needed!




## Number of Features and Sample Size

- If we have  $n$  samples in one dimension, we need  $n^m$  samples in  $m$  dimensions.
- $n^m$  increases exponentially with increase in number of input features.
- Adding more input features to improve performance works only if we can increase the number of input samples proportionately
- Adding more features without increasing sample size will worsen performance
- Practically, finding more data is very very hard!
- Dimensionality reduction is a more feasible approach: Data preprocessing, Principal component analysis



# How to Determine the Optimal Number of Features?

For a given sample size, there is an optimal number of features to use,



**Empirical approach:** Keep adding features until the error reaches a minimum before it begins to increase.

## Independent Validation

- Testing performance with training data is unreliable
- It may lead to overfitting
- This may be avoided by using a standard statistical procedure called *independent validation or cross-validation*
  - Randomize the dataset
  - Partition data into training and testing set
  - Further section the training data into 2 sets: training (or estimation) set and validation set
  - Use the training set to update the weights; Use the validation set to assess network performance
  - Use the test data to test how well the trained network generalizes
- This method of **cross-validation is called the holdout method**





# Using Cross-Validation for Model Selection

- Model selection is basically the problem of finding the optimal number of nodes (weights and biases) in the network
- We train the network and validate its performance for various network configurations; choose the one that minimizes the generalization error
- The optimal ratio of the sizes of the training set and validation set plays an important role in minimizing the generalization error.
- For  $N$  data samples,  $(1 - r)N$  samples are allocated to the training set and  $rN$  are allocated to the validation set.



## How do we specify $r$ ?

- low complexity problem → number of samples much larger than number of input features: performance is almost insensitive to  $r$
- High complexity problem: choice of  $r$  has a more pronounced effect on performance
- A single fixed value of  $r$  has been shown to work nearly optimally for a wide range of problems
- Kearns (1996) reported that  $r = 0.2$  is generally a sensible choice.

20% Validation samples, 80% Training samples

## Using Cross-Validation as an Early-Stopping Criterion

- Typically, as training progresses the mean square error decreases, and the network moves towards the minimum
- For good generalization, it is necessary to stop training at the 'right point' on the error curve.
- Stopping too early leads to poor generalization, stopping too late leads to overfitting.
- Cross-validation can be used to identify the onset of overfitting.




## Train then Validate Periodically

- Train the network for a period
- Periodically stop training, every 5 epochs for instance, present the validation input to the network and find the validation error
- Continue training and periodic training until the validation error reaches a minimum

This procedure is called the early-stopping method of training.


# Illustration of Early-Stopping



## Multifold Cross-Validation

- When there is a scarcity of data samples, we can use *multifold cross-validation*.
- Divide the dataset of  $N$  samples into  $K$  subsets
- Train the network with  $K-1$  subsets and validate with the  $K^{\text{th}}$  subset.
- Repeat until all subsets have been used for training and validation.
- Network performance is the average of the error over all trials.





This process is computationally expensive. The model has to be trained K times.

With a severely limited dataset, we may use *the leave-one-out* method. Train the network on N- 1 samples, test with the final sample. Network must be trained N times.

# Bias and Variance in a Neural Network

- Overfitting is one the most serious problems that arises in supervised learning networks.
- We use statistical techniques like independent validation to prevent overfitting.
- These methods are employed to balance the statistical bias and statistical variance when doing neural network learning to achieve smallest average generalization error.
- Statistical bias accounts for the degree of fitting the training data.
- The statistical variance accounts for the generalization ability of the network.
- Ideally we would like to minimize both statistical bias and statistical variance to improve performance.





# Bias-Variance Dilemma

- There is a tradeoff between bias and variance; cannot minimize one without increasing the other.
- A network that fits the data closely has low bias but a high variance. If we reduce variance this will lead to a decrease in the degree to which the data is fitted.
- To reduce both the statistical bias and variance we can add more data samples; not always possible.
- Statistical bias can be reduced by increasing the network size.
- Statistical variance can be decreased by pruning the neural network.

*Pruning refers to the removal of synaptic connections*

# Complexity Regularization

- The tradeoff between the reliability of data and goodness of the model can be represented as,  

$$R(w) = \varepsilon_{av}(w) + \lambda \varepsilon_c(w)$$
- The  $\varepsilon_{av}(w)$  term is the standard average error of the backpropagation algorithm, which depends on the both data and the network configuration.
- The  $\varepsilon_c(w)$  term is the complexity penalty measured in terms of the network weights.
- $\lambda$  is the regularization parameter representing the importance assigned to the complexity penalty.
- A small  $\lambda$  indicates unconstrained learning; A large  $\lambda$  indicates a constrained learning process.

*More on regularization theory in later lectures.*

# Weight-Decay Procedure

- A type of complexity regularization process
- The complexity penalty term is defined as,

$$\varepsilon_c(w) = \|w\|^2 = \sum_{i \in C_{total}}$$

where  $C_{total}$  refers to all synaptic weights in the network

- This procedure drives some network weights to zero; effectively disconnecting the weight.
  - Weights that have a significant impact on the performance have large values.
  - Weights that have little to no influence on performance are disconnected; also called excess weights.
  - Removal of network weights is called pruning; improves generalization.
- 



## Speed of Convergence

- Backprop is a generalization of LMS algorithm
- Learning parameter affects the speed of convergence
- This issue had led to the development of accelerated learning algorithms
- Heuristic improvement to standard backpropagation
  - Backprop with momentum (MATLAB: *traingdm*)
  - Variable learning rate gradient descent algorithms
- Use of numerical optimization techniques with backpropagation



# Variable Learning Rate Algorithms

## Adaptive learning rate backpropagation

- At each weight update step if the error increases by a given threshold, the new weights are discarded and a new (lower) learning rate is calculated.
- The new learning rate is decreased by a given fraction
- If the error decreases by the threshold amount, the learning rate is increased by a given fraction.
- MATLAB function: *traingda*

## Adaptive gradient descent with momentum

- MATLAB function: *traingdx*

# Accelerated Learning Algorithms

- Vector matrix form of backpropagation
- Conjugate gradient backpropagation
- Recursive least squares backpropagation
- Levenberg-Marquardt backpropagation



# Vector Matrix Form of Backpropagation

**Step 1:** Initialize the network synaptic weights to small random values

**Step 2:** Select a random input/output training sample, present it to the network and calculate the response of each layer simultaneously

**Step 3:** Use the output error to calculate the weight correction at each hidden layer,

- For the output layer  $j$ ,

$$\mathbf{D}_j = \mathbf{G}(\mathbf{v}_j)(\mathbf{d} - \mathbf{y}_j)$$

- For a hidden layer,

$$\mathbf{D}_{j-1} = \mathbf{G}(v_{j-1})\mathbf{D}_j\mathbf{W}_j$$

where,

- $\mathbf{G}(\mathbf{v}_j)$  is the gradient matrix at layer  $j$
- $\mathbf{W}_j$  is the weight at layer  $j$





- **Step 4:** Update the network weights according to,

$$\mathbf{W}_j(n+1) = \mathbf{W}_j(n) + \eta D_j(n) y_j(n)$$

where  $y_i(n) = x(n)$  for layer 1.

- **Step 5:** Continue steps 2 through 4, until the network reaches a desired level of performance





# Conjugate Gradient Backpropagation

- Conjugate gradient method is a numerical techniques for solving optimization problems
- Presents a compromise between simplicity of steepest descent and fast quadratic convergence of Newton's method
- The error function is approximated as a quadratic function

$$\mathcal{E}(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{Q} \mathbf{w} - \mathbf{b}^T \mathbf{w}$$

where  $\mathbf{w}$  is the vector of network weights and  $\mathbf{Q}$  is the Hessian matrix

$$\begin{bmatrix} \frac{\partial^2 \mathcal{E}}{\partial w_1^2} & \frac{\partial^2 \mathcal{E}}{\partial w_1 \partial w_2} & \cdots & \frac{\partial^2 \mathcal{E}}{\partial w_1 \partial w_m} \\ \frac{\partial^2 \mathcal{E}}{\partial w_2 \partial w_1} & \frac{\partial^2 \mathcal{E}}{\partial w_2^2} & \cdots & \frac{\partial^2 \mathcal{E}}{\partial w_2 \partial w_m} \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial^2 \mathcal{E}}{\partial w_m \partial w_1} & \frac{\partial^2 \mathcal{E}}{\partial w_m \partial w_2} & \cdots & \frac{\partial^2 \mathcal{E}}{\partial w_m^2} \end{bmatrix}$$

The dimension of the Hessian is equal to the total number of weights in the network.

- Computing the Hessian is impractical for even moderately large networks. Most conjugate gradient methods use alternative approaches for calculating the Hessian matrix.



**TABLE 4.3** Summary of the Nonlinear Conjugate-Gradient Algorithm for the Supervised Training of a Multilayer Perceptron

*Initialization*

Unless prior knowledge on the weight vector  $\mathbf{w}$  is available, choose the initial value  $\mathbf{w}(0)$  by using a procedure similar to that described for the back-propagation algorithm.

*Computation*

1. For  $\mathbf{w}(0)$ , use back propagation to compute the gradient vector  $\mathbf{g}(0)$ .
2. Set  $\mathbf{s}(0) = -\mathbf{g}(0)$ .
3. At time-step  $n$ , use a line search to find  $\eta(n)$  that minimizes  $\mathcal{E}_{av}(\eta)$  sufficiently, representing the cost function  $\mathcal{E}_{av}$  expressed as a function of  $\eta$  for fixed values of  $\mathbf{w}$  and  $\mathbf{s}$ .
4. Test to determine whether the Euclidean norm of the residual  $\mathbf{r}(n)$  has fallen below a specified value, that is, a small fraction of the initial value  $\|\mathbf{r}(0)\|$ .
5. Update the weight vector:

$$\mathbf{w}(n+1) = \mathbf{w}(n) + \eta(n)\mathbf{s}(n)$$

6. For  $\mathbf{w}(n+1)$ , use back propagation to compute the updated gradient vector  $\mathbf{g}(n+1)$ .
7. Set  $\mathbf{r}(n+1) = -\mathbf{g}(n+1)$ .
8. Use the Polak–Ribière method to calculate:

$$\beta(n+1) = \max\left\{\frac{\mathbf{r}^T(n+1)(\mathbf{r}(n+1) - \mathbf{r}(n))}{\mathbf{r}^T(n)\mathbf{r}(n)}, 0\right\}$$

9. Update the direction vector:

$$\mathbf{s}(n+1) = \mathbf{r}(n+1) + \beta(n+1)\mathbf{s}(n)$$

10. Set  $n = n + 1$ , and go back to step 3.

*Stopping criterion.* Terminate the algorithm when the condition

$$\|\mathbf{r}(n)\| \leq \varepsilon \|\mathbf{r}(0)\|$$

is satisfied, where  $\varepsilon$  is a prescribed small number.

## Comments on Conjugate-Gradient Algorithms

- Four variations of the Conjugate-Gradient method based on how the constant  $\beta$  is calculated.
  - Polak-Ribieri update (MATLAB: `trainscg`)
  - Fletcher-Reeves update (MATLAB: `trainscgf`)
  - Powell-Beale restarts (MATLAB: `trainscgb`)
  - Scaled conjugate gradient (MATLAB: `trainscgf`)
- The learning rate is calculated at each iteration in the algorithm

## Levenberg-Marquardt Backpropagation

- The Levenberg-Marquardt algorithm is a compromise between the Newton methods that converges rapidly near a minimum, and the gradient descent which is guaranteed to converge, but slowly.
- The basic update step of the Newton method is,

$$w(n+1) = w(n) - \mathbf{H}^{-1}(n)\mathbf{g}(n)$$

where  $\mathbf{H}$  is the Hessian matrix and  $\mathbf{g}$  is the error gradient.

- The LM algorithm was designed to improve the training speed by avoiding the calculation of the Hessian.



# Computation of Levenberg-Marquardt Weight Update Equation

Consider the approximation of the Hessian,

$$\mathbf{H} = \mathbf{J}^T \mathbf{J}$$

the error gradient is given by,

$$\mathbf{g} = \mathbf{J}^T \mathbf{e}$$

and  $\mathbf{J}$  is the Jacobian matrix that contains the first derivatives of the error function.

$$\begin{bmatrix} \frac{\partial e_1}{\partial w_1} & \frac{\partial e_1}{\partial w_2} & \dots & \frac{\partial e_1}{\partial w_m} \\ \frac{\partial e_2}{\partial w_1} & \frac{\partial e_2}{\partial w_2} & \dots & \frac{\partial e_2}{\partial w_m} \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial e_N}{\partial w_1} & \frac{\partial e_N}{\partial w_2} & \dots & \frac{\partial e_N}{\partial w_m} \end{bmatrix}$$

where  $m$  is the number of input variables and  $N$  is number of input samples



# Levenberg-Marquardt Weight Update

Levenberg-Marquardt weight update is given by

$$\underline{w(n + 1) = w(n) - [J^T J + \mu I]^{-1} e}$$

- Computing the Jacobian is much simpler than calculating the Hessian matrix, reducing computational complexity.

When  $\mu$  is 0, the weight update reverts to basic Newton's method.

When  $\mu$  is large the weight update becomes gradient descent with a small learning rate.

MATLAB function: *trainlm*





## Virtues and Limitations of Backpropagation

- Computational efficiency
- Connectionism
- Replicator mapping
- Function approximation
- Sensitivity analysis
- Robustness
- Convergence
- Local minima
- Scaling



[1] Ham and Kostanic, Principles of Neurocomputing for Science and Engineering, 2001.

[2] Kearns et. al., An Experimental and Theoretical Comparison of Model Selection Methods, Proc. Eighth ACM Conference on Computational Learning Theory.

[3] Simon Haykin, Neural Networks and Learning Machines, 3rd Ed., 2008.

All figures and content adapted from reference [3], unless specified otherwise.